

```

In [1]: 1 ### Importing libraries
2 import keras
3 from keras.datasets import mnist
4 from keras.layers import Dense,Dropout,Flatten,BatchNormalization
5 from keras.models import Sequential
6 from keras.layers import Conv2D,MaxPooling2D
7 from keras import backend as K

In [2]: 1 ## Loading dataset
2 (x_train,y_train),(x_test,y_test) = mnist.load_data()
3
4 ## network parameters
5 img_rows,img_cols = 28,28
6 batch_size = 128
7 n_epoch = 15
8 classes = 10
9
10 ##
11 if K.image_data_format() == 'channels_first':
12     x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
13     x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
14     input_shape = (1, img_rows, img_cols)
15 else:
16     x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
17     x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
18     input_shape = (img_rows, img_cols, 1)
19
20 ##Converting x_train,x_test to float
21 x_train = x_train.astype('float32')
22 x_test = x_test.astype('float32')
23
24 ## normalize it between 0-1
25 x_train /= 255
26 x_test /= 255
27
28
29 print('x_train shape:', x_train.shape)
30 print(x_train.shape[0], 'train samples')
31 print(x_test.shape[0], 'test samples')
32
33 ## one hot encode the target labels
34 y_train = keras.utils.to_categorical(y_train,classes)
35 y_test = keras.utils.to_categorical(y_test,classes)

```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)
 11493376/11490434 [=====] - 0s 0us/step
 x_train shape: (60000, 28, 28, 1)
 60000 train samples
 10000 test samples

Model 1 :

Architecture

Kernels : 3x3,5x5,2x2

Number of convnet layers : 3

Optimizer :Adam

Activation : Relu

Maxpooling layers : 3

Regularizations : Dropout,Batch Normalization after every two layers

```

In [3]: 1 model = Sequential()
2 ## create the architecture
3 model.add(Conv2D(32,activation='relu', kernel_size=(3, 3),input_shape = input_shape))
4 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
5 model.add(Conv2D(64, (5, 5), activation='relu',padding='same'))
6 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
7
8 model.add(BatchNormalization())
9 model.add(Dropout(0.35))
10
11 model.add(Conv2D(128, (2, 2), activation='relu',padding='same'))
12 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
13
14 model.add(BatchNormalization())
15 model.add(Dropout(0.35))
16
17 model.add(Flatten())
18 model.add(Dense(256, activation='relu'))
19 model.add(Dense(classes, activation='softmax'))
20
21 ## compile the model
22 model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam(),
23               metrics=['accuracy'])
24
25 ## train the model
26 history = model.fit(x_train,y_train,batch_size=batch_size,
27                     epochs=n_epoch,
28                     verbose=1,
29                     validation_data=(x_test, y_test))
30
31 score = model.evaluate(x_test, y_test, verbose=0)
32 print('Test loss:', score[0])
33 print('Test accuracy:', score[1])

```

```

Epoch 1/15
469/469 [=====] - 3s 7ms/step - loss: 0.1953 - accuracy: 0.9401 - val_loss: 0.1523 - val_accuracy: 0.9509
Epoch 2/15
469/469 [=====] - 3s 6ms/step - loss: 0.0637 - accuracy: 0.9801 - val_loss: 0.0351 - val_accuracy: 0.9887
Epoch 3/15
469/469 [=====] - 3s 7ms/step - loss: 0.0456 - accuracy: 0.9853 - val_loss: 0.0242 - val_accuracy: 0.9928
Epoch 4/15
469/469 [=====] - 3s 7ms/step - loss: 0.0376 - accuracy: 0.9881 - val_loss: 0.0326 - val_accuracy: 0.9895
Epoch 5/15
469/469 [=====] - 3s 7ms/step - loss: 0.0322 - accuracy: 0.9898 - val_loss: 0.0269 - val_accuracy: 0.9916
Epoch 6/15
469/469 [=====] - 3s 6ms/step - loss: 0.0290 - accuracy: 0.9909 - val_loss: 0.0320 - val_accuracy: 0.9897
Epoch 7/15
469/469 [=====] - 3s 7ms/step - loss: 0.0269 - accuracy: 0.9918 - val_loss: 0.0309 - val_accuracy: 0.9914
Epoch 8/15
469/469 [=====] - 3s 7ms/step - loss: 0.0235 - accuracy: 0.9923 - val_loss: 0.0233 - val_accuracy: 0.9928
Epoch 9/15
469/469 [=====] - 3s 7ms/step - loss: 0.0208 - accuracy: 0.9935 - val_loss: 0.0211 - val_accuracy: 0.9931
Epoch 10/15
469/469 [=====] - 3s 7ms/step - loss: 0.0196 - accuracy: 0.9939 - val_loss: 0.0365 - val_accuracy: 0.9890
Epoch 11/15
469/469 [=====] - 3s 7ms/step - loss: 0.0201 - accuracy: 0.9934 - val_loss: 0.0250 - val_accuracy: 0.9921
Epoch 12/15
469/469 [=====] - 3s 7ms/step - loss: 0.0189 - accuracy: 0.9938 - val_loss: 0.0379 - val_accuracy: 0.9910
Epoch 13/15
469/469 [=====] - 3s 6ms/step - loss: 0.0156 - accuracy: 0.9950 - val_loss: 0.0414 - val_accuracy: 0.9896

```

Epoch 14/15

469/469 [=====] - 3s 7ms/step - loss: 0.0159 - accuracy: 0.9951 - val_loss: 0.0286 - val_accuracy: 0.9927

Epoch 15/15

469/469 [=====] - 3s 6ms/step - loss: 0.0158 - accuracy: 0.9949 - val_loss: 0.0267 - val_accuracy: 0.9927


Test loss: 0.026714051142334938

Test accuracy: 0.9926999807357788

```

In [4]: 1  ## Lets plot the accuracy
        2  import matplotlib.pyplot as plt
        3  def plot_dynamic(x,y,y_,ax):
        4      ax.plot(x,y,label='Train_loss')
        5      ax.plot(x,y_,label='Test_loss')
        6      plt.legend()
        7      plt.grid()
        8      fig.canvas.draw()
        9
       10  x = list(range(n_epoch))
       11
       12  fig,ax = plt.subplots(1,1)
       13  ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
       14  vy = history.history['loss']
       15  ty = history.history['val_loss']
       16  plot_dynamic(x, vy, ty, ax)

```



Model 2 :

Architecture

Kernels : 3x3,5x5,2x2

Number of convnet layers : 5

Optimizer :Adam

Activation : Relu

Maxpooling layers : 5

Regularizations : Dropout,Batch Normalization

```

In [5]: 1 # from keras import backend as K
        2 # K.common.set_image_dim_ordering('tf')
        3
        4 model = Sequential()
        5 ## create the architecture
        6 model.add(Conv2D(16,activation='relu', kernel_size=(3, 3),input_shape = input_shape,padding='same'))
        7 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
        8 model.add(Conv2D(32, (5, 5), activation='relu',padding='same'))
        9 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
       10
       11
       12 model.add(BatchNormalization())
       13 model.add(Dropout(0.35))
       14
       15
       16 model.add(Conv2D(64, (2, 2), activation='relu',padding='same'))
       17 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
       18 model.add(Conv2D(128,(3, 3),activation='relu',padding='same'))
       19 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
       20
       21 model.add(BatchNormalization())
       22 model.add(Dropout(0.35))
       23
       24 model.add(Conv2D(256, (5, 5), activation='relu',padding='same'))
       25 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
       26
       27 model.add(Flatten())
       28 model.add(Dense(512, activation='relu'))
       29
       30 model.add(BatchNormalization())
       31 model.add(Dropout(0.35))
       32
       33 model.add(Dense(classes, activation='softmax'))
       34
       35 ## compile the model
       36 model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam(),
       37               metrics=['accuracy'])
       38
       39 ## train the model
       40 history = model.fit(x_train,y_train,batch_size=batch_size,\
       41                   epochs=n_epoch,\
       42                   verbose=1,\
       43                   validation_data=(x_test, y_test))
       44
       45 score = model.evaluate(x_test, y_test, verbose=0)
       46 print('Test loss:', score[0])
       47 print('Test accuracy:', score[1])

```

Epoch 1/15

469/469 [=====] - 4s 9ms/step - loss: 0.2421 - accuracy: 0.9243 - val_loss: 0.4402 - val_accuracy: 0.8528

Epoch 2/15

469/469 [=====] - 4s 8ms/step - loss: 0.0852 - accuracy: 0.9725 - val_loss: 0.0407 - val_accuracy: 0.9857

Epoch 3/15

469/469 [=====] - 4s 8ms/step - loss: 0.0625 - accuracy: 0.9798 - val_loss: 0.0366 - val_accuracy: 0.9892

Epoch 4/15

469/469 [=====] - 4s 8ms/step - loss: 0.0532 - accuracy: 0.9833 - val_loss: 0.0450 - val_accuracy: 0.9856

Epoch 5/15

469/469 [=====] - 4s 8ms/step - loss: 0.0446 - accuracy: 0.9860 - val_loss: 0.0344 - val_accuracy: 0.9897

Epoch 6/15

469/469 [=====] - 4s 8ms/step - loss: 0.0435 - accuracy: 0.9863 - val_loss: 0.0361 - val_accuracy: 0.9890

```


Epoch 7/15
469/469 [=====] - 4s 8ms/step - loss: 0.0375 - accuracy: 0.9881 - val_loss: 0.0333 - val_accuracy: 0.9902
Epoch 8/15
469/469 [=====] - 4s 8ms/step - loss: 0.0355 - accuracy: 0.9887 - val_loss: 0.0280 - val_accuracy: 0.9919
Epoch 9/15
469/469 [=====] - 4s 8ms/step - loss: 0.0327 - accuracy: 0.9896 - val_loss: 0.0319 - val_accuracy: 0.9916
Epoch 10/15
469/469 [=====] - 4s 8ms/step - loss: 0.0307 - accuracy: 0.9904 - val_loss: 0.0356 - val_accuracy: 0.9889
Epoch 11/15
469/469 [=====] - 4s 8ms/step - loss: 0.0296 - accuracy: 0.9905 - val_loss: 0.0331 - val_accuracy: 0.9910
Epoch 12/15
469/469 [=====] - 4s 8ms/step - loss: 0.0287 - accuracy: 0.9908 - val_loss: 0.0231 - val_accuracy: 0.9938
Epoch 13/15
469/469 [=====] - 4s 8ms/step - loss: 0.0266 - accuracy: 0.9913 - val_loss: 0.0266 - val_accuracy: 0.9921
Epoch 14/15
469/469 [=====] - 4s 8ms/step - loss: 0.0265 - accuracy: 0.9913 - val_loss: 0.0326 - val_accuracy: 0.9902
Epoch 15/15
469/469 [=====] - 4s 8ms/step - loss: 0.0244 - accuracy: 0.9923 - val_loss: 0.0232 - val_accuracy: 0.9930
Test loss: 0.02315954491496086
Test accuracy: 0.9929999709129333

```

```

In [6]: 1 x = list(range(n_epoch))
        2 fig,ax = plt.subplots(1,1)
        3 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
        4 vy = history.history['loss']
        5 ty = history.history['val_loss']
        6 plot_dynamic(x, vy, ty, ax)

```



Model 3 :

Architecture

Kernels : 3x3,5x5,2x2,7x7

Number of convnet layers : 7

Optimizer :Adam

Activation : Relu

Maxpooling layers : 7

Regularizations : Dropout,Batch Normalization after every two layers

```

In [7]: 1 model = Sequential()
2 ## create the architecture
3 model.add(Conv2D(16,activation='relu', kernel_size=(3, 3),input_shape = input_shape))
4 model.add(MaxPooling2D(pool_size=(2, 2)))
5 model.add(Conv2D(32, (5, 5), activation='relu',padding='same'))
6 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
7
8 model.add(BatchNormalization())
9 model.add(Dropout(0.35))
10
11 model.add(Conv2D(64, (2, 2), activation='relu',padding='same'))
12 model.add(MaxPooling2D(pool_size=(2, 2)))
13 model.add(Conv2D(128,activation='relu', kernel_size=(3, 3),padding='same'))
14 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
15
16 model.add(BatchNormalization())
17 model.add(Dropout(0.35))
18
19 model.add(Conv2D(256, (5, 5), activation='relu',padding='same'))
20 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
21 model.add(Conv2D(384, (7, 7), activation='relu',padding='same'))
22 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
23
24 model.add(BatchNormalization())
25 model.add(Dropout(0.35))
26
27 model.add(Conv2D(474, (5,5), activation='relu',padding='same'))
28 model.add(MaxPooling2D(pool_size=(2, 2),padding='same'))
29
30 model.add(Flatten())
31 model.add(Dense(512, activation='relu'))
32
33 model.add(BatchNormalization())
34 model.add(Dropout(0.35))
35 model.add(Dense(classes, activation='softmax'))
36
37 ## compile the model
38 model.compile(loss=keras.losses.categorical_crossentropy,optimizer=keras.optimizers.Adam(),
39               metrics=['accuracy'])
40
41 ## train the model
42 history = model.fit(x_train,y_train,batch_size=batch_size,
43                    epochs=n_epoch,
44                    verbose=1,
45                    validation_data=(x_test, y_test))
46
47 score = model.evaluate(x_test, y_test, verbose=0)
48 print('Test loss:', score[0])
49 print('Test accuracy:', score[1])

```

Epoch 1/15

469/469 [=====] - 10s 21ms/step - loss: 0.2895 - accuracy: 0.9096 - val_loss: 2.0209 - val_accuracy: 0.5539

Epoch 2/15

469/469 [=====] - 9s 20ms/step - loss: 0.1095 - accuracy: 0.9672 - val_loss: 0.0518 - val_accuracy: 0.9841

Epoch 3/15

469/469 [=====] - 9s 20ms/step - loss: 0.0797 - accuracy: 0.9761 - val_loss: 0.0550 - val_accuracy: 0.9844

Epoch 4/15

469/469 [=====] - 9s 20ms/step - loss: 0.0665 - accuracy: 0.9797 - val_loss: 0.0477 - val_accuracy: 0.9858

Epoch 5/15

469/469 [=====] - 9s 20ms/step - loss: 0.0611 - accuracy: 0.9820 - val_loss: 0.0362 - val_accuracy: 0.9887

```


Epoch 6/15
469/469 [=====] - 9s 20ms/step - loss: 0.0517 - accuracy: 0.9844 - val_loss: 0.0517 - val_accuracy: 0.9856
Epoch 7/15
469/469 [=====] - 10s 21ms/step - loss: 0.0454 - accuracy: 0.9869 - val_loss: 0.0388 - val_accuracy: 0.9892
Epoch 8/15
469/469 [=====] - 9s 20ms/step - loss: 0.0454 - accuracy: 0.9864 - val_loss: 0.0265 - val_accuracy: 0.9920
Epoch 9/15
469/469 [=====] - 9s 20ms/step - loss: 0.0422 - accuracy: 0.9875 - val_loss: 0.0300 - val_accuracy: 0.9910
Epoch 10/15
469/469 [=====] - 9s 20ms/step - loss: 0.0395 - accuracy: 0.9880 - val_loss: 0.0307 - val_accuracy: 0.9903
Epoch 11/15
469/469 [=====] - 9s 20ms/step - loss: 0.0385 - accuracy: 0.9884 - val_loss: 0.0377 - val_accuracy: 0.9901
Epoch 12/15
469/469 [=====] - 9s 20ms/step - loss: 0.0356 - accuracy: 0.9894 - val_loss: 0.0360 - val_accuracy: 0.9904
Epoch 13/15
469/469 [=====] - 9s 20ms/step - loss: 0.0324 - accuracy: 0.9902 - val_loss: 0.0285 - val_accuracy: 0.9915
Epoch 14/15
469/469 [=====] - 9s 20ms/step - loss: 0.0317 - accuracy: 0.9901 - val_loss: 0.0261 - val_accuracy: 0.9922
Epoch 15/15
469/469 [=====] - 9s 20ms/step - loss: 0.0306 - accuracy: 0.9911 - val_loss: 0.0248 - val_accuracy: 0.9921
Test loss: 0.024801477789878845
Test accuracy: 0.9921000003814697

```

```

In [8]: 1 x = list(range(n_epoch))
        2 fig,ax = plt.subplots(1,1)
        3 ax.set_xlabel('epoch')
        4 ax.set_ylabel('Categorical Crossentropy Loss')
        5 vy = history.history['loss']
        6 ty = history.history['val_loss']
        7 plot_dynamic(x, vy, ty, ax)

```



Observations:

- * As the number of convolution layer increases accuracy increases and loss reduces
- * As the number of convolution layers increases loss is reduces sharply at very early epochs.
- * Dropout and Batch normalization acts as regulaizers and avoids overfitting

